

## A MISSING PROOFS IN SECTIONS 4 AND 5

### A.1 Proof of Lemma 2

Lemma 2 tells the value ranges of  $n$  elements given their average value  $x$ , i.e., the  $i$ -th element must fall in  $[a_i, b_i]$  where  $a_i = \max\{x_l, \frac{x \cdot n - \sum_{j=i+1}^n x_h}{i}\}$  and  $b_i = \min\{x_h, \frac{x \cdot n - \sum_{j=1}^{i-1} x_l}{n-i+1}\}$ . It can be proved as follows.

**PROOF OF LEMMA 2.** Given a non-decreasing sequence  $\mathbf{x} = (x_1, \dots, x_n)$  with  $x_i \in [x_l, x_h]$  and  $\text{avg}(\mathbf{x}) = x$ , assume that there exists a element  $x_j$  in  $\mathbf{x}$  is out of  $[a_j, b_j]$ , i.e.,  $x_j \notin [a_j, b_j]$ . There are four cases to prove that such the  $x_j$  cannot exist.

Case 1:  $a_j = x_l$  and  $x_j < a_j$ . In this case,  $x_j$  is less than the lowest value of the domain, which violates the constraint  $x_j \in [x_l, x_h]$ .

Case 2:  $b_j = x_h$  and  $x_j > b_j$ . Similarly, the constraint  $x_j \in [x_l, x_h]$  is violated.

Case 3:  $a_j > x_l$  and  $x_j < a_j$ . One can construct a new sequence  $\hat{\mathbf{x}} = \{a_j, \dots, a_j, x_h, x_h, \dots, x_h\}$ , where the first  $j$  elements are  $a_j$  and the rest of the elements are equal to  $x_h$ . As  $\mathbf{x}$  is non-decreasing, we know that each element in  $\hat{\mathbf{x}}$  is no less than that in  $\mathbf{x}$ . Furthermore, there is  $\text{avg}(\hat{\mathbf{x}}) > \text{avg}(\mathbf{x})$  due to  $x_j < a_j$ . However, based on Eq. 4, there are  $a_j = \frac{x \cdot n - \sum_{k=j+1}^n x_h}{j}$  and  $\text{avg}(\hat{\mathbf{x}}) = x$ . It contradicts the fact that  $\text{avg}(\mathbf{x}) = x$ .

Case 4:  $b_j < x_h$  and  $x_j > b_j$ . Similarly, we can also construct a new sequence  $\hat{\mathbf{x}} = \{x_l, \dots, x_l, b_j, \dots, b_j\}$ , where the first  $(j-1)$  elements are  $x_l$  and other elements equal  $b_j$ . For the sequence  $\hat{\mathbf{x}}$ , its average value is less than  $\text{avg}(\mathbf{x})$ , since each element in  $\hat{\mathbf{x}}$  is no higher than that in  $\mathbf{x}$  and there is  $x_j > b_j$ . Besides, we get  $\text{avg}(\hat{\mathbf{x}}) = x$  based on Eq. 4, which is in conflict with  $\text{avg}(\mathbf{x}) = x$ .

Considering all cases together, the existence of  $x_j \notin [a_j, b_j]$  is a contradiction. As a result, the value range of any element  $x_i$  in the non-decreasing sequence  $\mathbf{x}$  is  $[a_i, b_i]$ . The proof is complete.  $\square$

### A.2 Proof of Theorem 2

We prove Theorem 2 case by case.

**PROPOSITION 4.** *Given the selection query  $Q_1$ : select  $A$  from  $R_j$  where predicates, the avg-count query  $Q_2$  and the count query  $Q_3$  with the same predicates,  $Q_1$  determines  $Q_2$  and  $Q_2$  determines  $Q_3$ . Further, we have that  $p_b(Q_3) \leq p_b(Q_2) \leq p_b(Q_1)$ .*

**PROOF.** Obviously,  $Q_1$  determines  $Q_2$  and  $Q_2$  determines  $Q_3$  due to the semantics of these queries. Moreover, for each tuple  $t \in R_j$ , there is  $E_t(Q_1) \subseteq E_t(Q_2) \subseteq E_t(Q_3)$ . We can prove this step by step.

First, when  $t$  satisfies the predicates of these three queries, its value ranges under the three queries  $Q_1, Q_2, Q_3$  are  $E_t(Q_1) = \{v|v \in S_j \text{ and } Q_1(v) = t.A\}$ ,  $E_t(Q_2) = \{v|v \in S_j \text{ and } Q_2(v) \in [a_i, b_i]\}$ , and  $E_t(Q_3) = \{v|v \in S_j \text{ and } Q_3(v) \neq \emptyset\}$ , respectively.  $[a_i, b_i]$  is the value range derived by Lemma 2. It is easy to find that  $E_t(Q_1) \subseteq E_t(Q_2) \subseteq E_t(Q_3)$  holds.

Second, the value ranges of  $t$  under the three queries are same as  $E_t = \{v|v \in S_j \text{ and } Q_1(v) = \emptyset\}$ . In this case,  $E_t(Q_1) \subseteq E_t(Q_2) \subseteq E_t(Q_3)$  also holds. Considering all cases together, there is  $E_t(Q_1) \subseteq E_t(Q_2) \subseteq E_t(Q_3)$  and the query price has  $p_b(Q_3) \leq p_b(Q_2) \leq p_b(Q_1)$ . It completes the proof.  $\square$

**PROPOSITION 5.** *Suppose there are two queries with the same predicates, i.e.,  $Q_1$ : select  $A$  from  $R_j$  where predicates and the*

*avg-count query  $Q_2$ : select avg( $A$ ), count( $*$ ) from  $R_j$  where predicates. When the number of qualified tuples of  $Q_1$  is  $n \leq 1$ , the two queries are equivalent and have the same price.*

**PROOF.** First, when  $n = 0$ , the two queries have empty answers and they are equivalent. In this case, all tuples in  $R_j$  under  $Q_1$  and  $Q_2$  have the same value range  $E_t = \{v|v \in S_j \text{ and } Q_1(v) = \emptyset\}$ . Thus, the prices of  $Q_1$  and  $Q_2$  are the same.

Second, when  $n = 1$ ,  $Q_1$  tells that, the value range of the one qualified tuple  $t$  is  $E_t(Q_1) = \{v|v \in S_j \text{ and } Q_1(v) = t.A\}$  and those of the rest of tuples are  $E_t = \{v|v \in S_j \text{ and } Q_1(v) = \emptyset\}$ . Under  $Q_2$ , the value range of the qualified tuple  $t$  on  $A$  can be computed based on Eq. 4, which is the same as  $E_t(Q_1)$ . Also, the rest of the tuples have the same value range  $E_t$  under  $Q_2$ . It means that,  $Q_1$  and  $Q_2$  have the same price. It completes the proof.  $\square$

**PROPOSITION 6.** *Suppose there are two queries with the same predicates, i.e.,  $Q_1$ : select  $A$  from  $R_j$  where predicates and the avg-count query  $Q_2$ : select avg( $A$ ), count( $*$ ) from  $R_j$  where predicates. When they both have the predicate ' $A = c$ ' and  $c$  is a constant, the two queries are equivalent and have the same price.*

**PROOF.** When  $Q_1$  and  $Q_2$  have the same predicates  $A = c$  where  $c$  is a constant, the query answer of  $Q_2$  is  $(c, n)$ , while  $Q_1$  returns a size- $n$  result  $(c, c, \dots, c)$ . In this case, the value ranges of  $n$  tuples satisfying the predicates are  $E_t(Q_1) = \{v|v \in S_j \text{ and } Q_1(v) = c\}$ . On the other hand, one can compute the value ranges of these tuples under  $Q_2$  based on Eq. 4. In this case, the  $a_i$  and  $b_i$  derived from Eq. 4 are  $c$ , since the minimum, the maximum, and the average of these tuples are  $c$ . Thus, these  $n$  tuples have the same value range, i.e.,  $E_t(Q_1)$ . Based on the same tuple information,  $Q_1$  and  $Q_2$  have the same price. The proof is complete.  $\square$

**PROPOSITION 7.** *Suppose there are two queries with the same predicates, i.e.,  $Q_1$ : select  $A$  from  $R_j$  where predicates and the avg-count query  $Q_2$ : select avg( $A$ ), count( $*$ ) from  $R_j$  where predicates. When all qualified tuples have the same value on the attribute  $A$ , the price of  $Q_2$  equals that of  $Q_1$ .*

**PROOF.** When  $Q_1$  and  $Q_2$  have the same predicates  $A = c$  where  $c$  is a constant, the query answer of  $Q_2$  is  $(c, n)$ , while  $Q_1$  returns a size- $n$  result  $(c, c, \dots, c)$ . In this case, the value ranges of  $n$  tuples satisfying the predicates are  $E_t(Q_1) = \{v|v \in S_j \text{ and } Q_1(v) = c\}$ . Specifically, the value ranges of the  $n$  qualified tuples under  $Q_2$  are computed based on Eq. 4. As these tuples have the same value on the attribute  $A$ , the maximum and minimum of the  $n$  tuples, i.e.,  $x_l$  and  $x_h$ , are the same as the average value. Hence, the  $a_i$  and  $b_i$  derived from Eq. 4 are also the same as the average, i.e., their real value. Thus, the value range of each qualified tuple under  $Q_1$  reveals its real value on the attribute  $A$ , which is the same as that under  $Q_1$ . In light of this, all tuples under  $Q_1$  and  $Q_2$  have the same value range and thus they have the same price. The proof is complete.  $\square$

**LEMMA 4.** *Suppose there are four queries with the same predicates, i.e.,  $Q_1$ : select  $A$  from  $R_j$  where same predicates, the avg-count query  $Q_2$ : select avg( $A$ ), count( $*$ ) from  $R_j$  where same predicates, the max/min query  $Q_3$ : select min( $A$ ) from  $R_j$  where same predicates, and the count query  $Q_4$ : select count( $*$ ) from  $R_j$  where same predicates and  $A < \bar{x}$  where  $\bar{x}$  is*

the average value. When all qualified tuples have the same value on the attribute  $A$ , we have that  $Q_2$  and  $Q_3$  determine  $Q_1$ , while  $Q_2$  and  $Q_4$  determine  $Q_1$  holds. Moreover, ARIA ensures no arbitrage, i.e.,  $p_b(Q_1) \leq p_b(Q_2) + p_b(Q_3)$  and  $p_b(Q_1) \leq p_b(Q_2) + p_b(Q_4)$ .

PROOF. Case 1.  $Q_2$  tells the average of  $n$  elements and  $Q_3$  tells their maximum/minimum. As these  $n$  elements have the same value, the maximum/minimum is the same as the average. In this case, data buyers can know based on  $Q_2$  and  $Q_3$  that, these elements must take the same value. Otherwise, the maximum/minimum cannot be the same as the average. In other words,  $Q_2$  and  $Q_3$  together tell the real values of these  $n$  elements, which are also revealed under  $Q_1$ . Hence,  $Q_2$  and  $Q_3$  determine  $Q_1$  in this case.

Case 2.  $Q_2$  tells the average of  $n$  elements and  $Q_4$  tells the number of elements that are less than the the average. As these  $n$  elements have the same value, the answer of  $Q_4$  is zero. Similarly, data buyers can know based on  $Q_2$  and  $Q_4$  that, these elements must take the same value. Otherwise, there must exist some elements less than the average. In this case,  $Q_2$  and  $Q_3$  actually tell the real values of these  $n$  elements, which are also revealed under  $Q_1$ . Thus, there is  $Q_2$  and  $Q_3$  determine  $Q_1$ .

Moreover, based on Proposition 7, there is  $p_b(Q_1) = p_b(Q_2)$ . Hence, there are  $p_b(Q_1) \leq p_b(Q_2) + p_b(Q_3)$  and  $p_b(Q_1) \leq p_b(Q_2) + p_b(Q_4)$ . The proof is complete.  $\square$

Before making effective proofs on the fourth case in Table 3, we first introduce the price of the limit query, i.e., the selection query with the limit  $k$  clause. The “limit  $k$ ” means that only  $k$  selections are returned. When the result size is  $k$ , i.e.,  $Q(D) = \{v_1, \dots, v_k\}$ , the buyer can know the value domains of  $k$  tuples, i.e.,  $E_t(Q) = \{v \in S_j \text{ and } Q'(v) = v_i\}$  where  $Q'$  removes the limit clause from  $Q$ . For other  $(N_j - k)$  tuples, they can either satisfy the query predicates or break the predicates, and thus their value domains are still  $S_j$ . When the result size is smaller than  $k$ , the limit query  $Q$  equals the query  $Q'$ , since the real result size is known. In this case, the value domain of each tuple can be derived as  $E_t(Q) = \{v|v \in S_j \text{ and } Q'(v) = Q'(t)\}$ . Based on the features of the prices of the limit query and avg-count query, we prove Lemma 5 based on Proposition 8. It means that, ARIA ensures no arbitrage where the original values of  $n$  elements are determined by the average and the  $(n - 1)$  elements.

PROPOSITION 8. Given any avg-count query  $Q$  with answer  $(\bar{x}, n)$  and  $n \geq 2$ , the total price of  $n$  qualified tuples under ARIA is no less than  $|S_j| - 1$ .

PROOF. There are at least two qualified tuples when  $n \geq 2$ . Let  $t_1$  (resp.  $t_2$ ) be the qualified tuple with minimum (resp. maximum) value on the aggregated attribute and  $[x_l, x_h]$  be the domain of this attribute. Based on Theorem 2, the value domain of  $t_1$  is  $E_1 = \{v|v \in_j \text{ and } Q'(v) \in [x_l, x]\}$ , while that of  $t_2$  is  $E_2 = \{v|v \in_j \text{ and } Q'(v) \in [x, x_h]\}$ .  $Q'$  is the non-aggregate version of  $Q$ . The total information gain of these two tuples is  $2 \cdot |S_j| - |E_1| - |E_2|$ . As  $E_1$  and  $E_2$  only overlaps at  $v$  with  $Q'(v) = x$ , there is  $|E_1| + |E_2| = |E| + 1$ , where  $E = \{v|v \in_j \text{ and } Q'(v) \in [x_l, x_h]\}$  and  $|E| \leq |S_j|$ . Therefore, the total information gain of all qualified tuples is no less than  $|S_j| - 1$ . The proof is complete.  $\square$

LEMMA 5. Given the queries  $Q_1$ : select  $A_k$  from  $R_j$  where predicates,  $Q_2$ : select  $A_k$ , count(\*) from  $R_j$  where predicates and  $Q_3$ : select  $A_k$  from  $R_j$  where predicates limit  $n - 1$ , i.e.,  $Q_1$  is determined by  $Q_2$  and  $Q_3$ , there is  $p_b(Q_1) \leq p_b(Q_2) + p_b(Q_3)$  where  $n$  is the number of qualified tuples in  $Q_1$ .

PROOF. When  $n = 0$  or  $n = 1$ , the avg-count query  $Q_2$  equals  $Q_1$ , and the derived tuple information of  $Q_2$  in Section 4.2 captures such phenomenon. In these cases, the price of  $Q_1$  is the same as that of  $Q_2$  and there is no arbitrage. When  $n$  is higher than 1, the price of  $Q_1$  considers the information gain of all tuples in  $R_j$ , i.e.,  $n$  qualified tuples and  $(N_j - n)$  unsatisfactory tuples. The information of unsatisfactory tuples is also revealed in  $Q_2$ , and their price under  $Q_1$  is the same as that in  $Q_2$ . For the first  $(n - 1)$  qualified tuples,  $Q_1$  and  $Q_3$  both reveal their information, and these tuples are priced as same. Then, for the last qualified tuple  $t$ , its price in  $Q_1$  is no higher than  $|S_j| - 1$ , since its value domain has at least one element, i.e., itself. Besides, Proposition 8 indicates that the total price of all qualified tuples in  $Q_2$  is no less than  $|S_j| - 1$ . Overall, the price of  $Q_1$  is no higher than the total price of  $Q_2$  and  $Q_3$  and there is no arbitrage caused by the  $(n - 1)$  freedom of the average value. The proof is complete.  $\square$

Based on the above propositions and lemma, we can prove Theorem 2 in the following.

PROOF OF THEOREM 2. Given a non-decreasing sequence  $\mathbf{x} = (x_1, \dots, x_n)$  with  $x_i \in [x_l, x_h]$  and  $\text{avg}(\mathbf{x}) = x$ , assume that there exists a element  $x_j$  in  $\mathbf{x}$  is out of  $[a_j, b_j]$ , i.e.,  $x_j \notin [a_j, b_j]$ . There are four cases to prove that such the  $x_j$  cannot exist.

Case 1: The selection query determines the avg-count query with the same predicates, while the avg-count query determines the count query. It is proved to ensure no arbitrage for such query determinacy by Proposition 4.

Case 2: The avg-count query can determine the selection query with the same predicates when the number of qualified tuples  $n$  is no higher than one or the predicates include the constant equation “ $A = c$ ”. The arbitrage-free property under this case is provided in Propositions 5 and 6.

Case 3: Proposition 7 ensures no arbitrage for the special case where all tuples satisfying the predicates of the avg-count query have the same value on the aggregated attribute.

Case 4: Lemma 4 achieves no arbitrage when the query determinacy comes from the degree of freedom of the average value.

As a result, ARIA is arbitrage-free to price avg-count queries considering all possible query determinacy in Table 3. The proof is complete.  $\square$

### A.3 Proof of Lemma 3

PROOF OF LEMMA 3. First, as  $Q$  is a selective query with the non-empty answer, there exists one tuple  $t_1$  satisfying the query predicates and another tuple  $t_2$  not. Then, for the tuple  $t_1$ , it can never take the value of  $t_2$ , i.e.,  $t_2 \notin E_{t_1}(Q)$  and  $|S_j| - |E_{t_1}(Q)| > 0$ . Hence, the price of  $t_1$  under  $Q$  is positive and the query price is higher than zero. The proof is complete.  $\square$

**Algorithm 5:** The ARIA-SAJ Algorithm

---

**Input:** the query  $Q$  on relations  $R_Q$ ; the sizes of all relations  $N_1, \dots, N_m$ ; the support sets  $S_1, \dots, S_m$

**Output:** the query price  $p$

- 1:  $Q' \leftarrow$  remove the max/min/avg/sum/count keyword, remove the group by clause, and add the ID column of each involved relation  $R_j \in R_Q$  into the selection clause of  $Q$
- 2: obtain  $Q'(D)$  and initialize the price  $p$  as zero
- 3: **foreach**  $R_j \in R_Q$  **do**
- 4:    $Q_j(D) \leftarrow$  extract the results of  $Q'(D)$  on relation  $R_j$ , remove repetitive results, and remove the ID column
- 5:    $Q_j(S_j) \leftarrow$  extract the results of  $Q'(D - R_j + S_j)$  on relation  $S_j$ , remove repetitive results, and remove the ID column
- 6:   **if** the aggregated attribute in  $Q$  is in  $R_j$  **then**
- 7:     **if** min/max in  $Q$  **then**
- 8:        $Q_j(D) \leftarrow$  group  $Q_j(D)$  and compute the maximum/minimum of each group
- 9:        $p_j \leftarrow$  ARIA-MM( $Q_j(D)$ ,  $Q_j(S_j)$ ,  $N_j$ ,  $|S_j|$ )
- 10:     **else**
- 11:        $p_j \leftarrow$  ARIA-AVG( $Q_j(D)$ ,  $Q_j(S_j)$ ,  $N_j$ ,  $|S_j|$ )
- 12:     **else**
- 13:        $p_j \leftarrow$  ARIA-Base( $Q_j(D)$ ,  $Q_j(S_j)$ ,  $N_j$ ,  $|S_j|$ )
- 14:      $p \leftarrow p + p_j$
- 15: **return**  $p$

---

**B MISSING ALGORITHM IN SECTION 5**

We describe the missing ARIA-SAJ algorithm in Section 5 to price the select-aggregate-join (SAJ) queries. For simple count queries, the SAJ query can be rewritten as the SJ query, which replaces the count(\*) with 1. For other types of aggregations, each SAJ query is decomposed into one select-aggregate (SA) query and multiple selection queries.

Algorithm 5 depicts the procedure of pricing SAJ queries in ARIA. It takes the SAJ query  $Q$  on multiple relations  $R_Q$ , the sizes of all relations (i.e.,  $N_1, \dots, N_m$ ), and all support sets (i.e.,  $S_1, \dots, S_m$ ) as inputs. It outputs the query price  $p$ . It first constructs the query  $Q'$  based on  $Q$ , which removes the special keywords and clauses (e.g., max, min, avg, sum, count, group by) and adds the ID column of each relation  $R_j \in R_Q$  into the selection clause of  $Q$  (line 1). Next, it executes  $Q'(D)$  to derive  $Q_j(D)$  and initializes the query price  $p$  as zero (line 2). Then, ARIA-SPJ starts to derive the price of each single-relation query  $Q_j$  (lines 3-10). It first derives  $Q_j(D)$  by manually extracting the results of  $Q'(D)$  on  $R_j$ , removing the repetitive results due to the join operation, and dropping the ID column. Similarly, ARIA-SPJ can obtain the query answer  $Q_j(S_j)$  (line 5).  $Q'$  is executed on the relation  $S_j$  and other relations (except for  $R_j$ ) in  $D$ , while the query answer is denoted as  $Q'(D - R_j + S_j)$ .

Given  $Q_j(D)$  and  $Q_j(S_j)$ , ARIA-SAJ starts to compute the price of each single-relation query. For the relation where the aggregate attribute is located, ARIA-SAJ needs to consider the price of the SA query (lines 7-11). In particular, ARIA-SAJ manually computes the max/min value of each group based on  $Q_j(D)$  and employs ARIA-MM to compute the price of max/min query. When  $Q$  is the avg-count/sum-count query, ARIA-SAJ directly utilizes the ARIA-AVG algorithm in line 11. Otherwise, ARIA-SAJ employs ARIA-Base to compute the price  $p_j$  of the non-aggregate query  $Q_j$  (line 13). The derived price  $p_j$  is accumulated to the query price  $p$ . When all involved relations are considered, the query price  $p$  of the SAJ query  $Q$  is returned (line 15).

The complexity of pricing the max/min queries on joined tables is  $\sum_{R_j \in R_Q} |S_j|$ , since ARIA-Base and ARIA-MM both have linear time complexity. Moreover, it requires  $O(\sum_{R_j \in R_Q} |S_j| \cdot \log \sum_{R_j \in R_Q} |S_j|)$  complexity of pricing avg-count/sum-count queries, the log-linear complexity  $O(|S_j| \cdot \log |S_j|)$  of pricing avg-count queries.

Received 17 July 2024; revised 12 March 2024; accepted 5 June 2024